

Open-Source AI Installation Guide (Windows)

Option 1: LM Studio (Easiest - No Command Line)

What it is

A standalone desktop app with built-in model management, chat UI, and local API server.

Steps

1. Go to <https://lmstudio.ai>
2. Download the Windows installer
3. Run the installer (defaults are fine)
4. Launch LM Studio
5. Go to the **Discover** tab and search for a model (e.g., "Llama 3.1", "Qwen 2.5 Coder")
6. Click **Download**
7. Once downloaded, go to the **Chat** tab and start chatting

Pros

- No command line needed
 - Built-in model browser and chat UI
 - Can run a local API server (OpenAI-compatible)
-

Option 2: Ollama (Command Line + Optional Web UI)

What it is

A lightweight CLI tool for running open-source LLMs. Pairs with Open WebUI for a browser-based chat interface.

Step 1: Install Ollama

1. Go to <https://ollama.com/download/windows>
2. Download and run the installer
3. Follow the prompts (defaults are fine)

Step 2: Download a Model

Open **PowerShell** or **Command Prompt** and run:

```
ollama pull llama3.1
```

Other recommended models:

```
ollama pull qwen2.5-coder
ollama pull deepseek-r1:8b
ollama pull phi4-mini
```

Step 3: Chat via Command Line

```
ollama run llama3.1
```

Type your messages and press Enter. Type `/bye` to exit.

Step 4 (Optional): Add Web UI

Install Python 3.11+ from <https://www.python.org/downloads/> then run:

```
pip install open-webui
open-webui serve
```

Open <http://localhost:8080> in your browser.

Option 3: Jan (Desktop App, Open Source)

What it is

An open-source desktop app similar to LM Studio.

Steps

1. Go to <https://jan.ai>
 2. Download the Windows installer
 3. Run the installer
 4. Launch Jan
 5. Download models from within the app
 6. Start chatting
-

Recommended Models by Use Case

Use Case	Model	RAM Needed
General chat	llama3.1	8 GB
Coding	qwen2.5-coder	8 GB
Reasoning	deepseek-r1:8b	8 GB
Powerful	llama3.1:70b	48 GB
Lightweight	phi4-mini	4 GB

Hardware Requirements

- **Minimum:** 8 GB RAM (can run 7-8B parameter models)
 - **Recommended:** 16 GB RAM (can run 14B parameter models)
 - **NVIDIA GPU:** Dramatically speeds up inference (used automatically by Ollama and LM Studio)
 - **No GPU:** Still works using CPU, just slower
-

Security Notes

- All data stays on your machine - nothing is sent to the cloud
- No accounts or API keys required

- Download models only from official sources (Ollama library, LM Studio catalog, HuggingFace official repos)
 - Verify model publishers (Meta, Mistral, Qwen, DeepSeek) to avoid tampered models
-

Connecting to VS Code (Optional)

With Ollama:

1. Install the **Continue** extension in VS Code
2. Configure it to use `http://localhost:11434` (Ollama's default)

With LM Studio:

1. Enable the local server in LM Studio (Developer tab)
 2. Install the **Continue** extension in VS Code
 3. Point it to `http://localhost:1234/v1`
-

Generated on 2026-05-08